

# Media Streaming

CloudSuite1.0 Benchmark Suite  
Copyright (c) 2011, Parallel Systems Architecture Lab, EPFL

All rights reserved.

The media streaming benchmark consists of two main components: A client and a server. The client component emulates real world clients; sending requests to stress a streaming server. The length and quality of videos requested by the client can be controlled by the client interface with the default setup closely following real world observations. The client system is used to send requests to a media streaming server that supports the RTSP protocol, like the Darwin Streaming Server used in our case.

## Prerequisite Software Components

1. [Darwin Streaming Server](#)
2. [Faban](#) workload generator kit
3. Stream driver, RTSP client code, and operation scripts (with the benchmark package)
4. [cURL](#) RTSP library

[Download](#) the Media Streaming Benchmark

## Installing the client

1. The client is built on top of the Faban workload generator. Therefore, the first step to install the client is to download Faban. You can obtain Faban from the Media Streaming Benchmark [package](#) or Faban's [website](#).

untar Faban in the desired directory. Call this directory: *faban-streaming*.

2. Copy the *streaming* directory from the downloaded package to the, just created, *faban-streaming* directory
3. Make sure that the Java JDK is installed and the JAVA\_HOME environment variable is properly set.
4. Start the Faban master: From Faban main directory (*faban-streaming*):

```
./master/bin/startup.sh
```

5. Download the streaming benchmark directory (called *streaming*), unpack it and copy it to the Faban main directory (*faban-streaming*).

The streaming directory contains the streaming benchmark driver and the RTSP client code, in addition to the necessary scripts that will ease the deployment and run of the benchmark.

6. In the streaming directory, make the following necessary modifications to *build.properties*.
  - o Change *faban.home* to the path of the main Faban directory (e.g., */home/username/faban-streaming*).
  - o Make sure *faban.url* is set to *http://localhost:9980/*.
7. From the streaming directory, run: *sh scripts/new\_depoly.sh*.

This script will compile the stream driver using the default values of the benchmark. Note that this step requires the Apache Ant tool installed.

8. Prepare the RTSP client

- The code is provided in the same streaming directory, *faban-streaming/streaming/rtsplibclientfinal.c*, you need to download the cURL version that works on your platform.
- unpack the downloaded cURL distribution.
- To install cURL, in the curl directory:
  - `mkdir curlinst`
  - `./configure --prefix=/path/to/curldirectory/curlinst/`
  - `make`
  - `make install`
- Now, compile the RTSP client code

```
gcc -I /path/to/curl/curlinst/include -L /path/to/curl/curlinst/lib/ -lcurl
rtsplibclientfinal.c -o rtsplibclient.o
```

- Make a preliminary test:

```
./rtsplibclient.o
```

If the installation was done correctly, you should receive the following message:  
"ERROR: enter a valid URL" simply because it expects a parameter indicating where to find the video.

## Installing the server

Here, we provide the steps necessary to install the Darwin Streaming Server on a Linux machine. Our test was carried out on Ubuntu10.10, but the instructions should apply to other distributions without problems.

1. Download and unpack the Darwin Streaming Server (on a different machine).
2. Download the following two patches:
  - [Patch one](#)
  - [Patch two](#)
3. Apply the patches:

```
patch -p1 < dss-6.0.3.patch
patch -p1 < dss-hh-20080728-1.patch
```

4. Replace the Install script with this [one](#) (if you are not using the benchmark package) and then change its permissions:

```
chmod +x Install
```

5. The Darwin streaming server requires a qtss group and qtss user in that group to start:

```
sudo addgroup qtss
sudo adduser --system --no-create-home --ingroup qtss qtss
```

6. Build the server: `./Buildit`
7. Install the server: `./Install`
8. Now the server is ready. To start the server:  
`/usr/local/sbin/DarwinStreamingServer -dDS 1`  
The server should print the main statistics every second!

Note: All the configuration parameters of the server exist in the `/etc/streaming/streamingserver.xml` file. You can specify parameters such as the maximum number of concurrent users and maximum bandwidth to use ...etc.

## Generating the data set

The streaming benchmark dataset consists of two main folders. The first, called `streamingVideos_10`, is a 1.6 GB folder and includes all the following video categories and lengths:

**Duration:** 1 minute to 10 minutes at the increments of 1 minute

**Rates:**

Dial-up: 42 Kbps

Low: 102 Kbps

Medium: 290 Kbps

Hi: 790 Kbps

LAN: 1500 Kbps

The second directory, called `stressVideos`, consists of three videos: 5-minute Dialup, 5-minute Medium and 5-minute Hi. The intent of these videos is to emulate real world situations where a small subset of the videos gets accessed most of the time.

To generate a large data set, you need to replicate the `streamingVideos_10` folder the number of times until the total dataset is sufficient for you. The replication process should use the same name but changing the number at end (`streamingVideos_11`, `streamingVideos_12`, etc.) and the numbers should be contiguous with no gaps.

In the case of the Darwin streaming server, the data set should be under the `/usr/local/movies` directory in the server machine. Otherwise, you can point the server to use your directory of choice by editing `/etc/streaming/streamingserver.xml` and modifying the corresponding option.

The benchmark comes with a simple script to help the process of dataset generation. To generate the dataset under the `/usr/local/movies` directory:

- Copy the *streamingVideos\_10* directory to the */usr/local/movies* directory.
- `sh copy_streaming.sh X`, where X is the number of directories you want to generate.

Note: This process can be time consuming depending on the size of the generated data set.

## Running the Benchmark

Before running the benchmark you should do the following:

1. `export JAVA_HOME=/usr/lib/jvm/java-6-sun` (or point to the JDK on your machine)
2. Make sure to start the master process: From the *faban-streaming* directory:  
`sh master/bin/startup.sh`

### From the command line:

The benchmark comes with a set of scripts to ease running the benchmark. The scripts necessary to run the benchmark are located under the *faban-streaming/streaming* (client directory) directory under the directory *scripts*.

The main script in this folder is *run-test.sh*. This script generates 10 Java processes, called agents, each of which will generate an equal number of requests to the server machine. For example, if the total number of requests is 100, each agent will generate 10 requests.

The *run-test.sh* script by itself reads the default configuration found in */path/to/faban-streaming/streaming/deploy/run.xml* and starts generating requests based on the settings in that file.

To use this script, first you need to set the correct paths as follows, you need to be in */path/to/faban-streaming/streaming* and:

- `vim scripts/run-test.sh` and replace all the */home/username/* to point to the path to *faban-streaming*
- `sh scripts/run-test.sh`

The following is a description of the main parameters in the *run.xml* file that you may need to modify to customize your run:

- Host IP address: specifies the address of the machine where the client agents will be running. This is set to *localhost* by default.
  - The scale: defines the number of clients simultaneously simulated.
  - The ramp up, steady state and ramp down time.
- Please note that the benchmark is designed to ramp up the clients based on the number of

such emulated clients. Therefore, the ramp up time is not critical but should be greater than or equal to 5 seconds. The important parameter for a successful run is the ramp down time. This value should be set to  $2 * \text{duration of longest video}$  you are willing to use.

- The number of agents: how many agent processes to control the total number of clients. Please note that this number affects the speed of preparing the client threads. Along with this parameter, you should change the *run-test.sh* script by including separate commands for all the agents. For example if you want to use 15 agents instead of 10, you should add 10 extra lines in the *run-test.sh*.
- Whether to start the agents concurrently or not: selecting yes starts all the agents at the same time, each creating the number of clients assigned to that agent.
- The mix of videos (duration and bitrate) you want your clients to use. For example, to stream only short-duration and high bitrate videos you need to set the *GetShortHi* parameter to 100 and the rest to zero. If you want the requested videos to be 30% short and high and 70% medium duration and low bitrate, you should specify the *GetShortHi* to 30 and the *GetMediumLow* to 70 and the rest to zeros, and so on.
- IP Address of the server
- The port number the streaming server uses to listen to client requests. Darwin Streaming Server listens to 4 ports by default, one of them is 554 (specified by default). You can check and set other ports by looking at (and editing) */etc/streaming/streamingserver.xml* in the server machine. You can also force the server to listen to a certain port by using the *-p* option in the Linux command line when you start the server. For example:  
*/usr/local/sbin/DarwinStreamingServer -dDS 1 -p 3330 to listen to port 3330*
- Data set scale: corresponds to the number of video directories that were generated (as part of the data set). If you want to generate ~3 GB of video data set you need to use two directories. To ensure that your requests actually request videos from both directories, you should specify this number as 2.

Please note that this number has no effect on the requests to *stressVideos*. For example, if you specify in the request mix that you need 100% of requests to *stressHi*, the dataset scale becomes irrelevant and will not be used, however, it will still be used for all other requests if their percentage in the mix is greater than 0.

- Fix duration: is an optional parameter that is provided to give more control over the user run. It limits the duration of each requested video to this number. For example, if you want to use a mix of videos with different durations and bit rates but want to test only 1 minute of each video, you should specify this number as 1. A negative number means ignoring this parameter (by default it is -1 meaning that it is ignored).
- Protocol code: is the directory where the RTSP client resides. The default case is under */path/to/faban-streaming/streaming*.

To facilitate the configuration process, the benchmark comes with few scripts to allow the rapid change and run of basic configurations. You can use the *start-run.sh* script in the */path/to/faban-*

*streaming/streaming/scripts* as follows:

From */path/to/faban-streaming/streaming*

*sh scripts/start-run.sh <NUM-CLIENTS> <MIX> <SERVER-IP> <DATASET-SCALE>*

where:

- NUM-CLIENTS: is the same as scale (number of simultaneous clients).
- SERVER-IP: is the IP address of the server.
- DATASET-SCALE: the same as the data set scale described earlier.
- MIX:
  - dist: a selected mix of durations and bitrates.
  - shorthi = 100% short duration high bitrate.Similarly, you can use: shortmedium, shortlow, longlow, longmedium, longhi, mediumlow, mediummedium, medium hi, stresshi, stressmedium, stresslow

### **From the GUI:**

1. To run the benchmark from the the gui, start by deploying the benchmark: open *http://localhost:9980/deploy* in a browser.

Username: deployer and password is: adminadmin

Then select *\$FABAN\_HOME/benchmarks/streaming/stream.jar* and click on *deploy*. You should see a message that says that the benchmark was deployed successfully!

2. Open *http://localhost:9980/* in a browser. You will see three main tabs: The first defines the java parameters, the second defines the main client parameters and the third defines the server parameters. Please refer to the previous section for a definition of the parameters.
3. After selecting the proper configuration, start the benchmark run by clicking the "ok" button.

### **Quality of Service**

The media streaming benchmark provides two types of clients; each comes with a different way of checking QoS.

1. Regular Client
2. QoS Client

QoS can be estimated in both clients, but in a different way. The benchmark uses the Regular Client by default.

The QoS client tries to accurately check QoS by recording the packet information that arrives for each client request every second. The benchmark driver then calculates the Quality seen by each individual client and provides the result of the benchmark run. To use the QoS client, you need root privileges, and do the following:

- From `/path/to/faban-streaming/streaming` replace `src/sample/streamdriver/StreamDriver.java` with `src-qos/sample/streamdriver/StreamDriver.java` (you may want to make a copy of the original to use it later)
- In the directory `scripts`, modify all the paths in the `.sh` files to point to the correct `/path/to/faban-streaming/`
- `sh scripts/new_deploy.sh`
- `sh scripts/clean_experiment.sh`
- `sh scripts/sample_iptables.sh` & (this script will sample the `iptables` each second for the duration of the test, you should kill it at the end).
- Run the client following the same procedure as described in the previous section)
- At the end of the run, the client will have the summary of the results in the `./results/AggregateResults` file that will explicitly mention whether the client has succeeded or failed the run.

The Regular Client only stresses the server with client requests but doesn't calculate the QoS provided by the server. QoS can be estimated by other means, such as the server stats, *wireshark* and other supporting scripts provided by the media streaming benchmark package. Because of the simplified nature of this client, its memory utilization is lower than the detailed one, and thus can be used if the benchmark user wants to stress the server in a faster way to roughly estimate Quality while using fewer client machines. This can be especially useful for the case where a limited number of machines can be used for the client or while porting the workload to a simulation environment.

The media streaming package provides the `analyze` script that can provide an estimation of the quality of service based in the average bit rate received for the overall number of videos requested in addition to another estimation based on each individual client request. Measuring the quality using this method requires the benchmark user to take a dump of the ethernet traffic using a tool like `tcpdump`. For example: `tcpdump -i any udp -w ./dumpdir/dump.pcap`

Then analyzing it using a tool like `capinfos`. For example: `capinfos -L ./dumpdir/dump.pcap`

where this command will get a report that includes the average bitrate received, from which the user can estimate whether the quality is acceptable or not. Taking an accurate measurement using this method requires considering the number of packets dropped with the average size of the packet to make the estimation.

The benchmark comes with a script called `analyze.sh` under the `scripts` folder. This script does the job of estimating the average bitrate on a per-client basis, and provides a summary of the requests that were served with a good quality. The benchmark user needs to provide the number of clients used in experiment and the bitrate expected for the videos (the experiment must use a single uniform bitrate). For example:

```
cd ./dumpdir
sh analyze.sh 1500 290
```